

Accelerating Permutation Testing in Voxel-wise Analysis through Subspace Tracking: A new plugin for SnPM

Felipe Gutierrez-Barragan, Vamsi K. Ithapu, Chris Hinrichs, Camille Maumet, Sterling C. Johnson, Thomas E. Nichols, and Vikas Singh

<http://felipegb94.github.io/RapidPT/>

The paper

- **Problem:** Permutation testing is computationally expensive in voxel-wise analysis.
- **Solution:** A model that leverages the structure of the permutation testing matrix to reduce the computation runtime.
- **Contributions:**
 - *Theoretical framework + Algorithm*
 - *RapidPT:* A MATLAB toolbox for fast and robust permutation testing.
 - 2x – 38x faster than state of the art (SnPM toolbox) and 20x – 1000x faster than a naïve permutation testing implementation.
 - *SnPM Plugin:* RapidPT is incorporated into a widely used neuroimaging toolbox. Not a common contribution of new algorithms.

Schedule

1. Background: Voxel-wise analysis in neuroimaging studies, multiple hypothesis testing, p-values and thresholds, controlling FWER.
2. Permutation Testing Procedure
3. The permutation testing matrix, \mathbf{T}
4. RapidPT Algorithm
5. Evaluations
 1. Accuracy
 2. Runtime gains
6. SnPM Plugin
7. Discussion

Background: Problem

- Consider a study with n subjects from two groups (ex: sick vs. healthy, resting state vs. non resting state)
- For each subject we have a v dimensional measurement vector (**voxels**, genes, region of interest, etc)
- **Question:** Does the data display any interesting activity? If so, where is that “interesting” activity? How certain are we?
- **Goal:** Classify each voxel as either active or not active and associate a probability (α) to how certain we are of our classification of each voxel.

Background: Hypothesis Testing

- **Solution:** Hypothesis Testing does exactly what we want.
 - Calculate the probability that your claim/hypothesis is true.

Null Hypothesis vs. Alternate Hypothesis


Inactive vs. Active

H_0 vs. H_a

- Basic idea:
 - Assume voxels from both groups come from the same distribution (null hypothesis).
 - Test if they do. If they do accept null, if they don't reject null.

Background: Hypothesis Testing - Procedure

1. Choose appropriate test statistic (t-test, mean difference, etc).
2. State H_0 and H_a
3. Construct the null distribution for the test statistic. This is the distribution of the statistic given that H_0 is true. This can be done analytically in some cases.
4. State **alpha**
5. Compute the test statistic with the given data.
6. Calculate the probability of observing such a statistic given that H_0 is true (**p-value**).
7. Accept or reject H_0 .



Hypothesis Testing - Types of Error

Actual Situation “Truth”		
		Decision
H₀ True	H₀ False	
Do Not Reject H₀	Correct Decision $1 - \alpha$	Incorrect Decision Type II Error β
Rejct H₀	Incorrect Decision Type I Error α	Correct Decision $1 - \beta$

$$\alpha = P(\text{Type I Error}) \quad \beta = P(\text{Type II Error})$$

Multiple Testing Problem

- **FWER:** Probability of making at least Type I Error (False Positive).

$$P(\text{Making at least 1 error in } m \text{ tests}) = 1 - (1 - \alpha)$$

$$P(\text{Making at least 1 error in } m \text{ tests}) = 1 - (1 - \alpha)^m$$

- Common Alpha = 0.05
- For $m = 100 \rightarrow \text{FWER} = 0.99$.
- In neuroimaging $m > 100,000$ tests. One for every voxel.
- Just by chance we will reject the null MANY times.
- Hence, we want to somehow control for the FWER and keep it under a certain probability α_0 .

Controlling FWER

- **Parametric Methods**

- Assumptions about the data and its distribution.
- **Bonferroni** – Conservative. Simply set $\alpha_0 = \alpha/v$ ($v = \text{number of tests}$).
- **Random Fields Theory** – Estimate number of activation areas, effectively reducing the number of tests.

- **Nonparametric/Resampling Methods**

- **Permutation Testing** – Empirically estimate the null distribution of the test statistics.
 - Exact control of the FWER.

Permutation Testing

- **Idea:** If the two groups do not differ, then I can permute the group/class labels and end up with approximately same set of test statistics.
- **Procedure:**
 1. Re-label images (permute the labels of the images).
 2. Compute test statistic for each voxel
 3. Repeat N times.
- After the procedure is done we will have the exact null distribution for each voxel, and we can proceed from step 4 of hypothesis testing .

Permutation Testing – Example, Single Test


	<u>Control</u>				<u>Drug</u>			
Expression	9	12	14	17	18	21	23	26
Average	13				22			

8 choose 4 = 70 possible permutations

Rearrangement of data									
Random Assignment	Control				Drug				Difference in Averages
1	9	12	14	17	18	21	23	26	9.0
2	9	12	14	18	17	21	23	26	8.5
3	9	12	14	21	17	18	23	26	7.0
4	9	12	14	23	17	18	21	26	6.0
5	9	12	14	26	17	18	21	23	4.5
6	9	12	17	18	14	21	23	26	7.0
7	9	12	17	21	14	18	23	26	5.5
8	9	12	17	23	14	18	21	26	4.5
9	9	12	17	26	14	18	21	23	3.0
10	9	12	18	21	14	17	23	26	5.0
11	9	12	18	23	14	17	21	26	4.0
12	9	12	18	26	14	17	21	23	2.5
13	9	12	21	23	14	17	18	26	2.5
14	9	12	21	26	14	17	18	23	1.0
15	9	12	23	26	14	17	18	21	0.0
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
69	18	21	23	26	9	12	14	17	-8.5
70	18	21	23	26	9	12	14	17	-9.0


14

Distribution of Difference between Treatment Means
Assuming No Treatment Effect



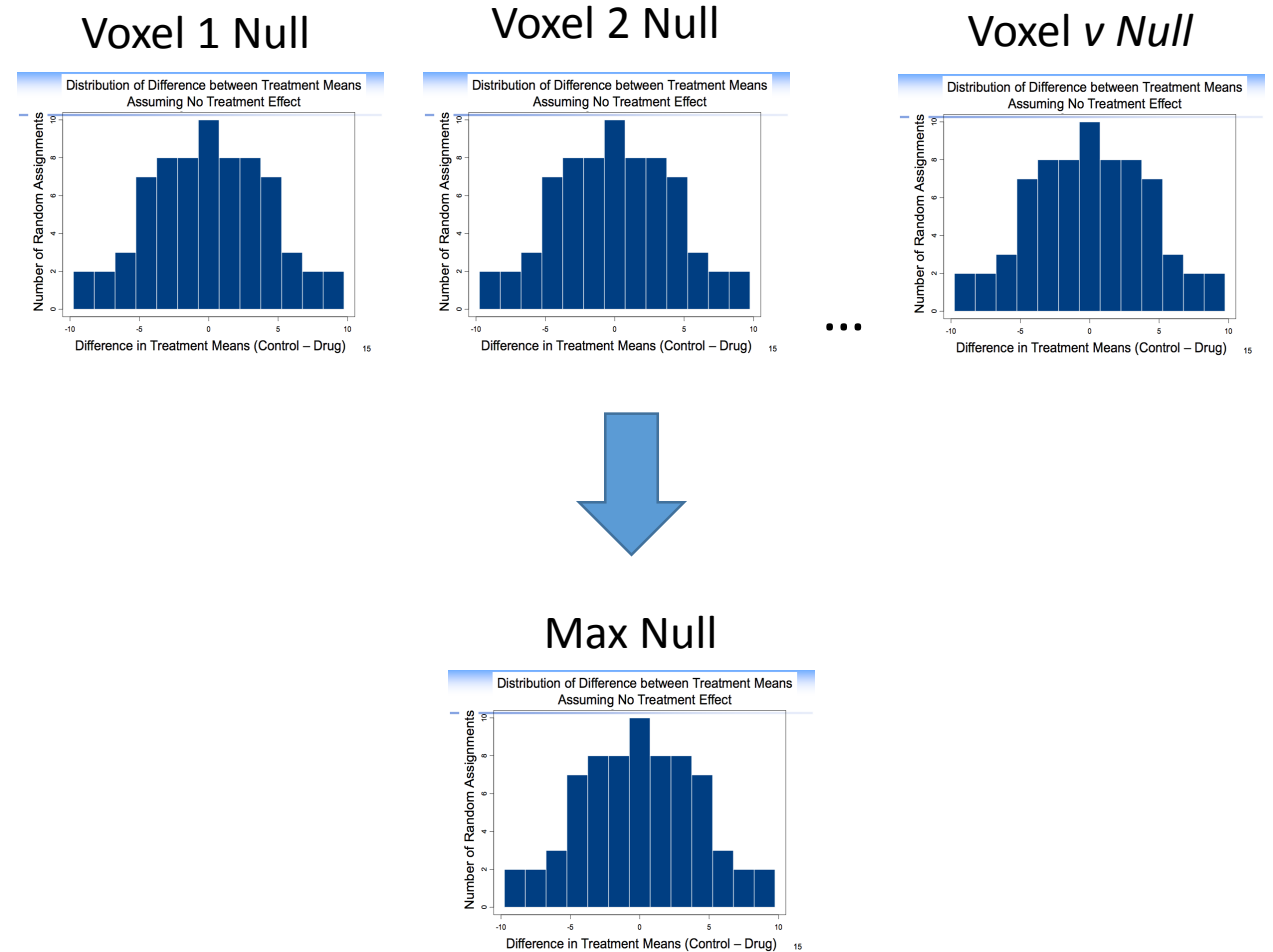
Permutation Testing Matrix

- For each voxel we will have one null distribution associated to it.
- The permutation testing matrix $\mathbf{T} = \mathbf{v} * \mathbf{L}$ matrix of test statistics.
- *Issue:* If $v \sim 100,00$, $L \sim 10,000$ the matrix is around 7.5 Gigabytes.




Maxnull Distribution

- We don't have to keep track of all distributions.
- Simple *keep track of the maximum test statistic across voxels for each permutation.*
- Construct the maxnull distribution



We still have computational issues...




- The possible number of relabelings becomes huge as the number of subjects increases.
 - $30 \text{ choose } 15 = 155,117,520$
- So we compute only a subset of them.
- For each relabeling of the data we have to compute v test statistics and find the max!
- More permutations = we can compute lower p-values.
- Embarrassingly parallel problem, BUT not all neuroscientists have access to expensive hardware/supercomputers, and most of them prefer to work on their laptops.


Idea: Look at the structure of T

- T is “highly structured” – A combination of low-rank signal and high-rank residual
- Example: MRI data 100 healthy vs. non-healthy. $v = 1,000$, $L = 2,000$

Permutation Testing Matrix P



Singular values of P



Core of RapidPT

- Many columns in T look similar to other columns as well as many rows look similar to other rows.
 - *Rank of T is constrained by the number of subjects.*
- If we compute a small number of entries of T we should be able to estimate the rest of it.
- Mathematically,

$$T = UW + S$$

- U is the low-rank basis of T .
 - W is the coefficient matrix
 - S is a high-rank random residual (noise).
- How many entries? In our experiments, subsampling <1% was enough

RapidPT Algorithm

- Divide the process into training and recovery.
- Training:
 - Calculate a few exact permutations $\sim N$
 - Estimate low-rank basis through subspace tracking
 - Estimate the residual

$$S = T_{exact} - UW$$
- Recovery:
 - For each permutation Calculate a subset of the test statistics (eta) for a column of T .
 - Recover the rest through matrix completion.
 - Get max test statistic

Algorithm 2 The RAPIDPT algorithm for permutation testing.

Input: $X^1, X^2, r, \eta, L, \ell, \text{stat}$

Output: \hat{T}, h^L

$X = [X^1; X^2], n = n_1 + n_2$

TRAINING

$U \leftarrow \text{RAND. ORTH.}, W_{ex} = [\emptyset]$

for $i \in 1, \dots, \ell$ **do**

$j_1, \dots, j_n \sim \text{PERMUTE}[1, n]$

$\tilde{X}^1 \leftarrow X[:, j_1, \dots, j_{n_1}]$

$\tilde{X}^2 \leftarrow X[:, j_{n_1+1}, \dots, j_n]$

$T_{ex}[:, i] \leftarrow \text{test}(\tilde{X}^1, \tilde{X}^2)$

$k_1, \dots, k_{[\eta v]} \sim \text{UNIF}[1, v]$

$\tilde{T} \leftarrow T_{ex}[k_1, \dots, k_{[\eta v]}, i]$

$U, W_{ex}[:, i] \leftarrow \text{SUBSPACE-TRACKING}(r)$

end for

$\sigma \leftarrow \text{STANDARD DEVIATION}\{T_{ex} - UW_{ex}\}_\Omega$

$\mu \leftarrow \sup_i \text{MAX}\{T_{ex}[:, i] - UW_{ex}[:, i]\}$

for $i \in 1, \dots, \ell$ **do**

$\hat{T}[:, i] \leftarrow T[:, i]$

end for

RECOVERY

for $i \in \ell + 1, \dots, L$ **do**

$k_1, \dots, k_{[\eta v]} \sim \text{UNIF}[1, v]$

$j_1, \dots, j_n \sim \text{PERMUTE}[1, n]$

$\tilde{X}^1 \leftarrow X[k_1, \dots, k_{[\eta v]}, j_1, \dots, j_{n_1}]$

$\tilde{X}^2 \leftarrow X[k_1, \dots, k_{[\eta v]}, j_{n_1+1}, \dots, j_n]$

$\tilde{T} \leftarrow \text{test}(\tilde{X}^1, \tilde{X}^2)$

$W[:, i] \leftarrow \text{COMPLETE}(U, \tilde{T}, k_1, \dots, k_{[\eta v]})$

$s \leftarrow \text{i.i.d.}\mathcal{N}^v(0, \sigma^2)$

$\hat{T}[:, i] \leftarrow UW[:, i] + s$

end for

for $i \in 1, \dots, L$ **do**

if $i \leq \ell$ **then**

$m_i \leftarrow \text{MAX}(\hat{T}[:, i])$

else

$m_i \leftarrow \text{MAX}(\hat{T}[:, i]) + \mu$

end if

end for

$h^L \leftarrow \text{HISTOGRAM}(m_1, \dots, m_L)$

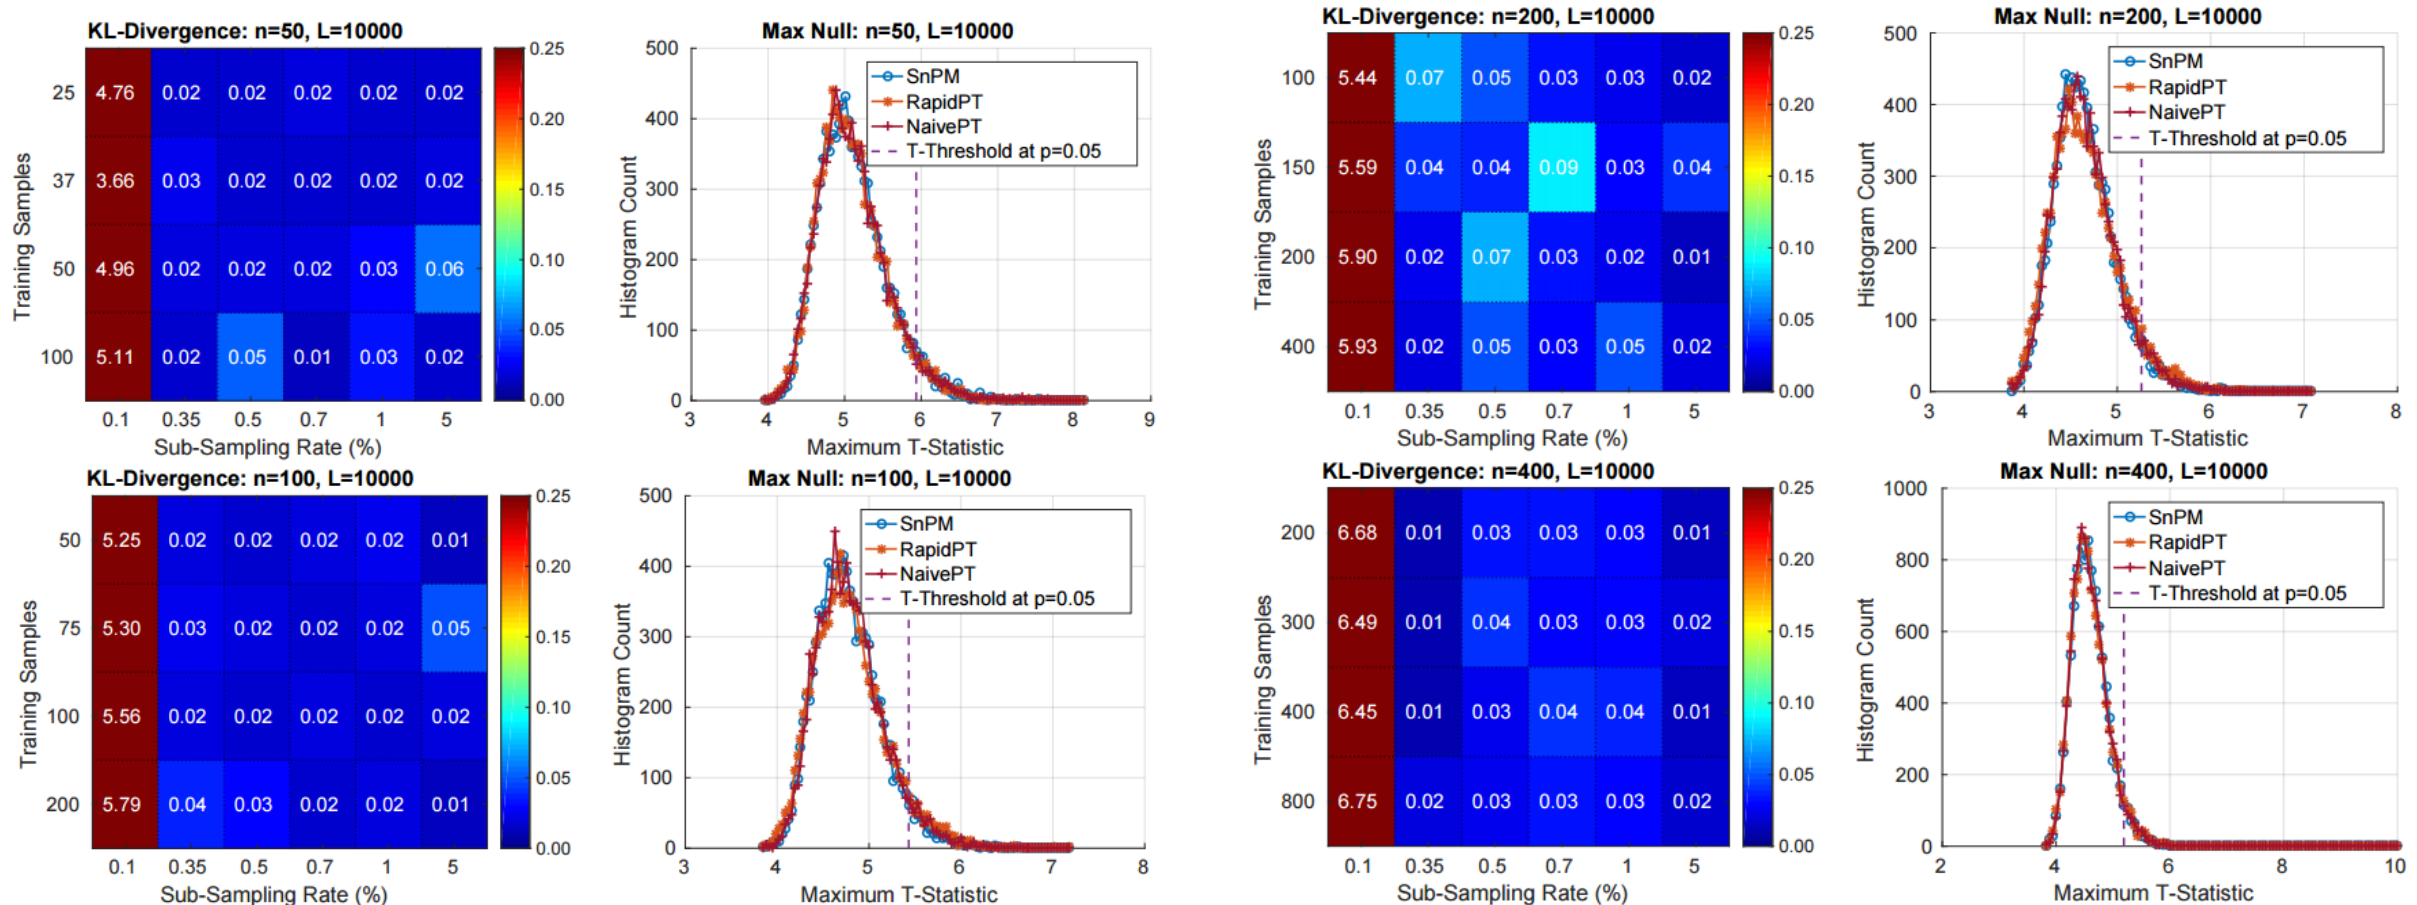
Evaluations Setup

- **Data:** T1 MRIs from ADNI2 are used.
 - 601 subjects (259 AD and 342 CN)
 - SPM preprocessing is applied.
 - GM images with approx. 500,000+ voxels are extracted.
 - Multiple combinations of dataset sizes
- **Experiments:** Can we recover the Maxnull distribution?
 - Stability of hyperparameters – Sub-sampling rates, and training samples
 - Computational Speedups (RapidPT vs. SnPM, RapidPT vs. NaivePT)

Recovered MaxNull Distribution Accuracy

KL-Divergence: Measure of the difference between two distributions

Datasets: 50, 100, 200 and 400 subjects



What Sub-Sampling Rate to Choose?

Low-rank matrix completion says that around $\sim r \log(d)$ entries are needed, where r is the column space rank and d is the ambient dimension.

$$\eta v \geq n \log(v)$$

$$\eta \geq \frac{n \log(v)}{v}$$

Empirically the above inequality worked well. However, we can sub-sample at lower rates

In our experiments, the number of subjects and number of voxels for each dataset were:

Number of Subjects and Number of Voxels: (n, v)			
(50,547783)	(100,558295)	(200,568086)	(400,574640)

Table 1: Number of subjects and number of voxels per subject on each dataset.

Therefore, the corresponding estimate of the minimum sub-sampling rate would be:

Minimum sub-sampling rate estimate: η_{min}			
$\approx 0.1206\%$	$\approx 0.2370\%$	$\approx 0.4665\%$	$\approx 0.9231\%$

Table 2: Number of subjects and number of voxels per subject on each dataset.



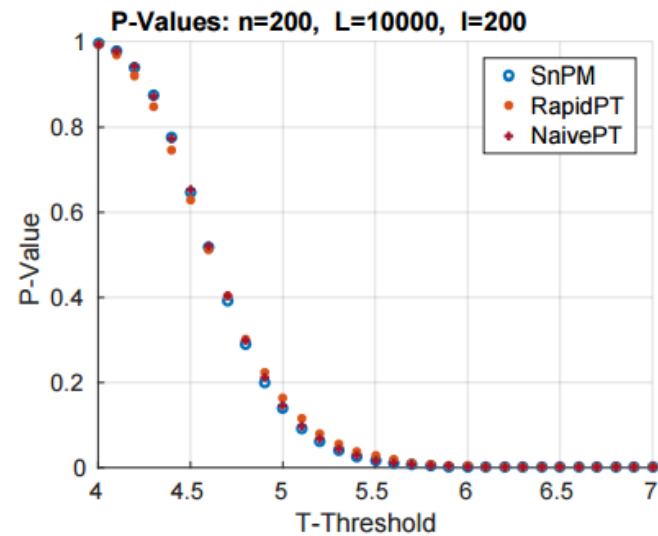
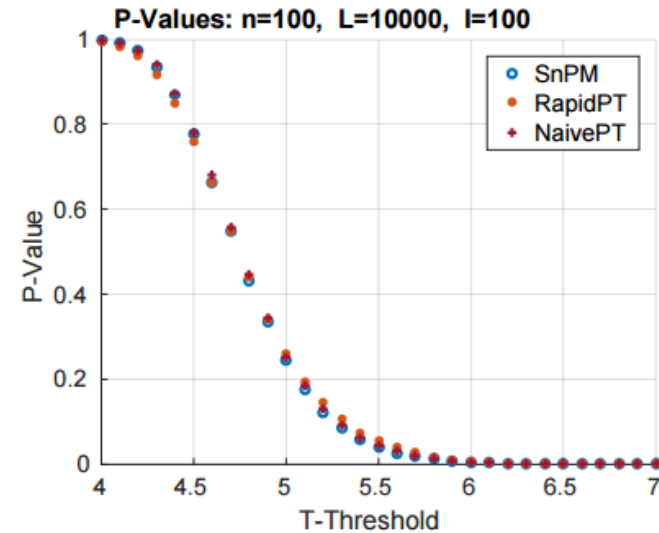
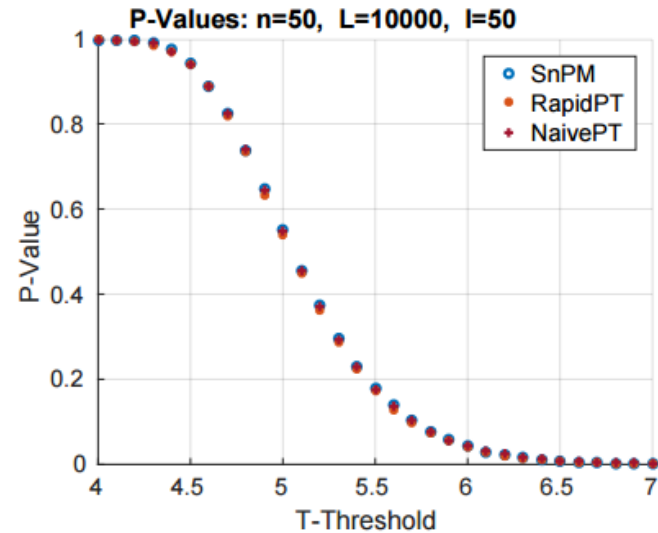
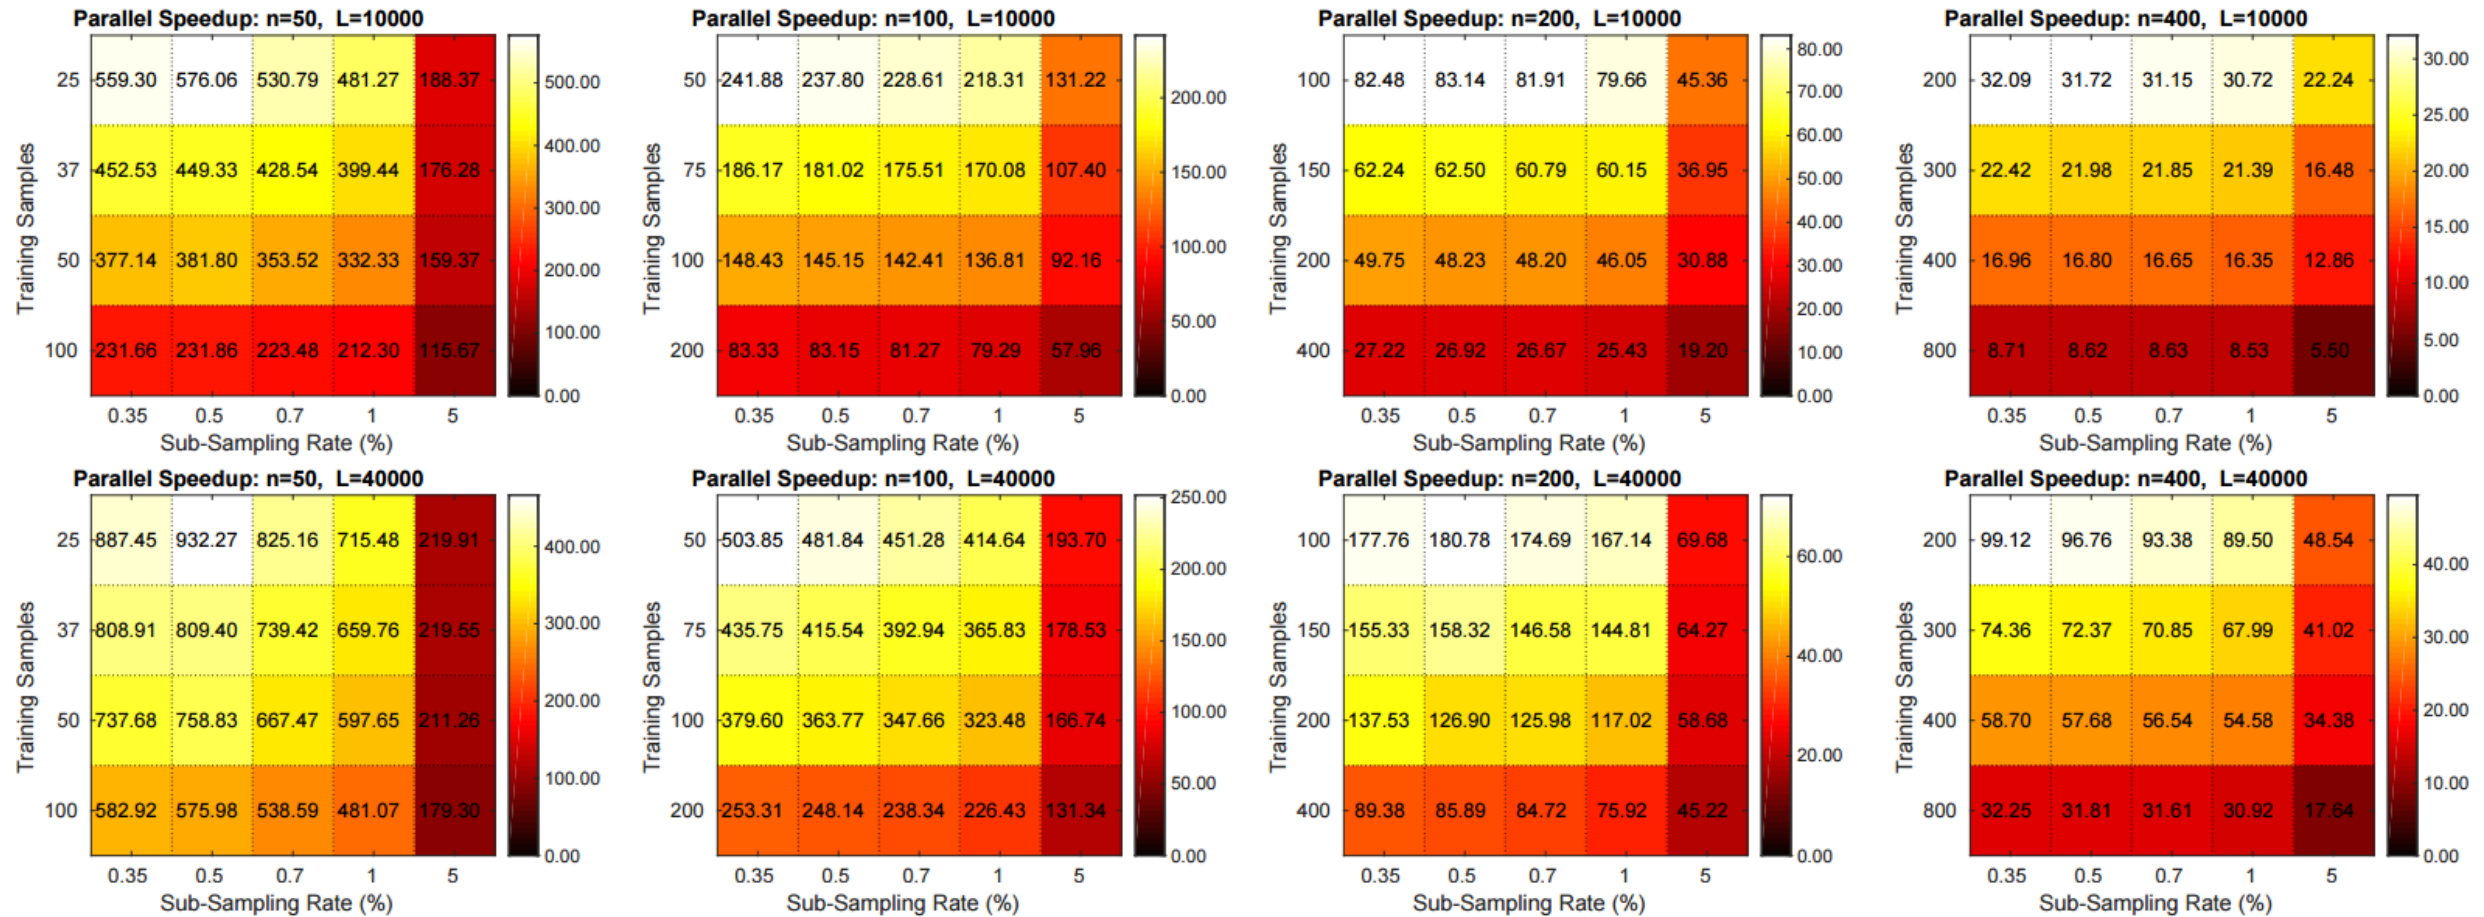


Figure 11: An experiment showing the effect of the sub-sampling rate (η) on the KL-Divergence.

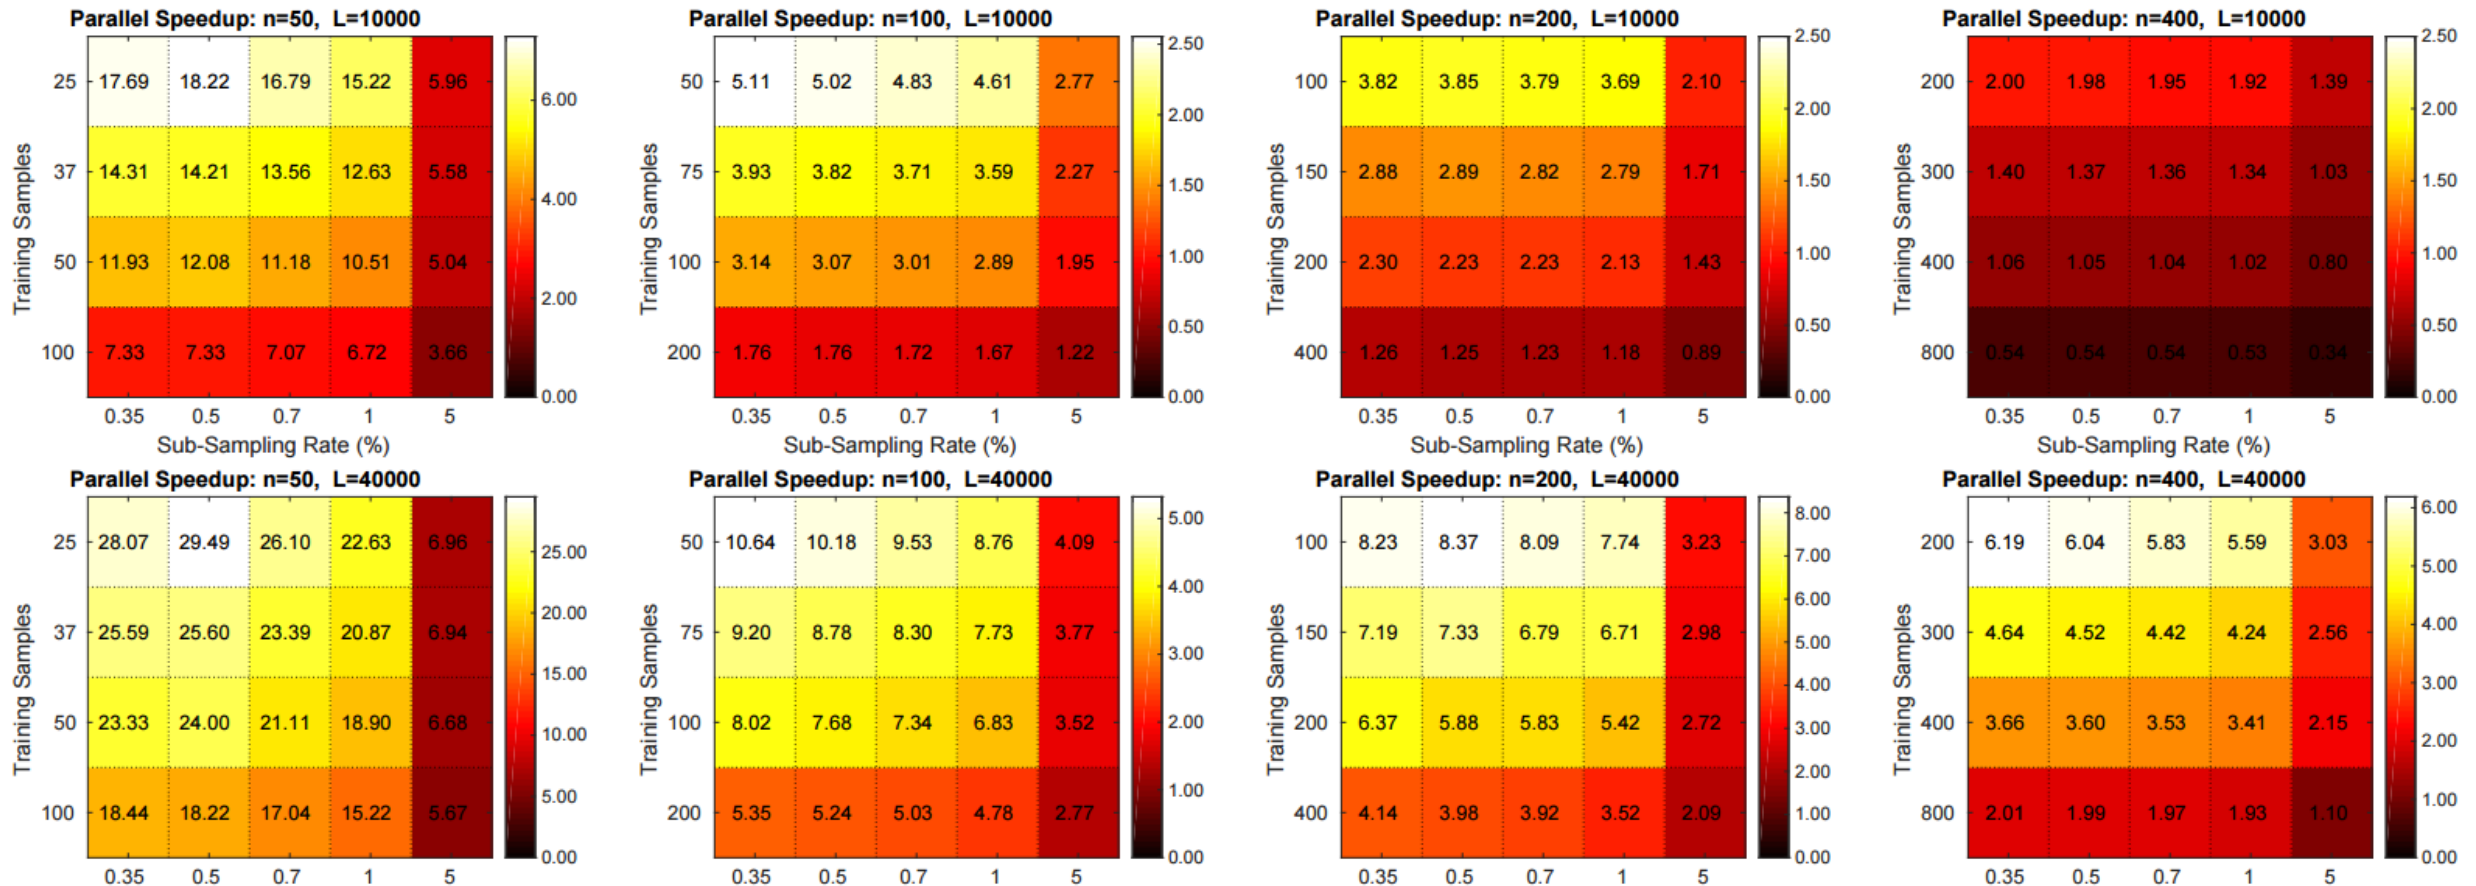
Calculated T-Threshold for a Given P-Value



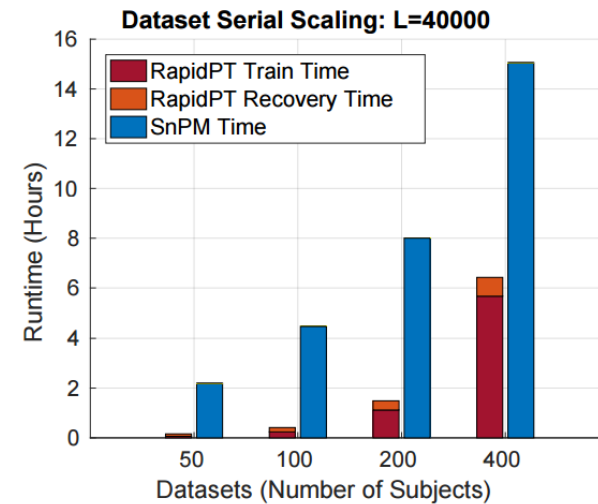



Speedup RapidPT vs. NaivePT







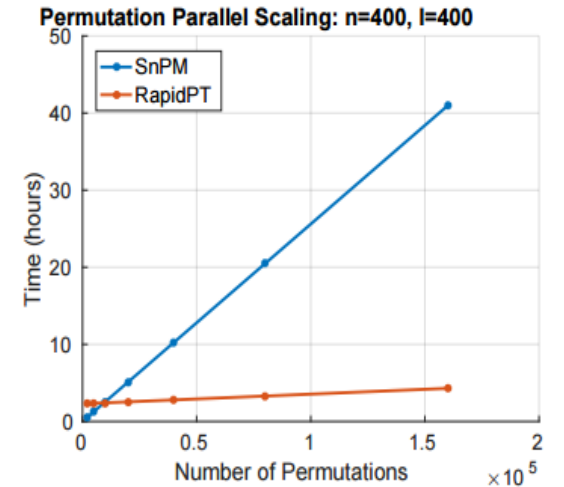



Speedup RapidPT vs. SnPM



Timing RapidPT vs. SnPM



Scaling the number of permutations

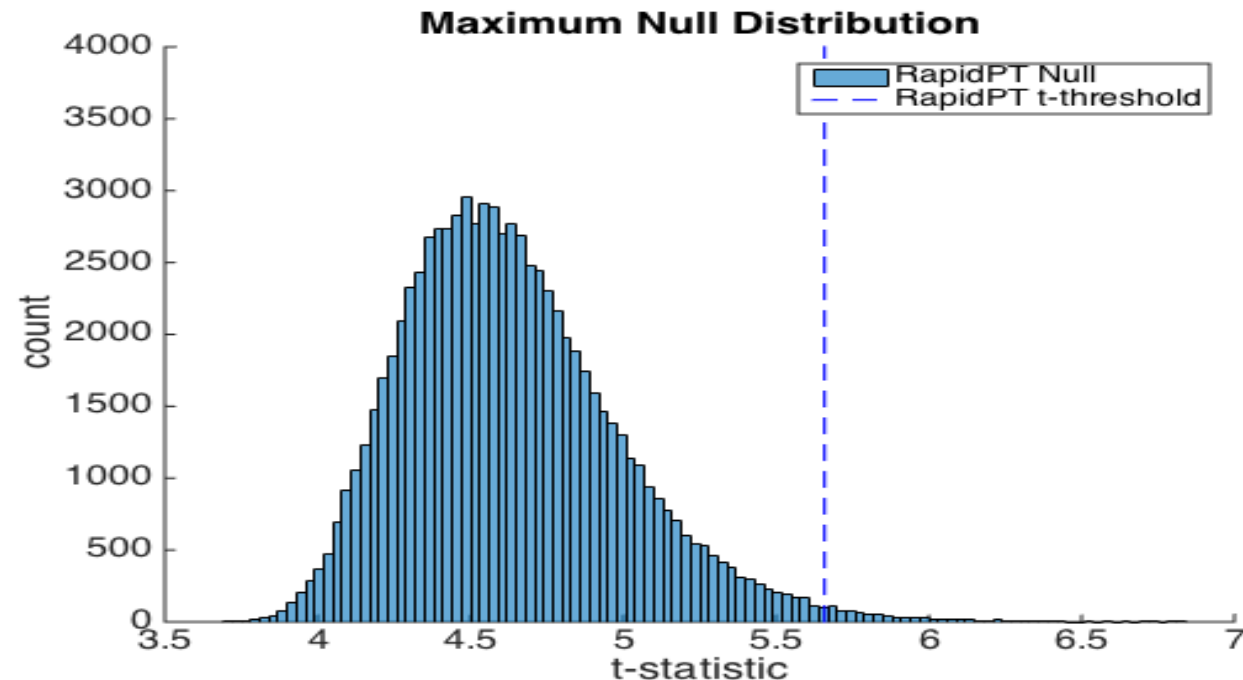


Discussion

- When to use RapidPT?
 - Many permutations. If you are doing $> 10,000$ permutations then RapidPT will highly benefit.
- When not to use it?
 - Permutation testing on small datasets take only a few seconds...
 - Large datasets > 200 subjects, and $< 10,000$ permutations

RapidPT – Postprocess Example

```
% Get the outputs struct you obtained from RapidPT
load('~\PermTest/outputs/TwoSample_ADRC_200_200_400/rapidpt/outputs_80000_0
alpha = 0.01; % Significance level of 1 percent
tThresh_RapidPT = prctile(outputs.maxT, 100 - (100*alpha));
% Get the data
load('~\PermTest\data/ADRC/TwoSample/ADRC_400_200_200.mat');
[h,p,ci,stats]=ttest2(Data(1:200,:),Data(201:400,:),0.05,'both','unequal');
SampleMaxT = max(stats.tstat);
```



Acknowledgements and Website

- Vikas Singh
- Vamsi Ithapu
- Chris Hinrichs, Camille Maumet, Sterling C. Johnson, Thomas E. Nichols
- ADNI and ADRC
- Repository and project website:
 - <https://github.com/felipegb94/RapidPT> (Repository)
 - <http://felipegb94.github.io/RapidPT/> (Website)