

CHRONO::HPC





DISTRIBUTED MEMORY FLUID-SOLID INTERACTION SIMULATIONS

Felipe Gutierrez, Arman Pazouki, and Dan Negrut
University of Wisconsin – Madison


Support: Rapid Innovation Fund, U.S. Army TARDEC


ASME IDETC/CIE 2016 :: **Software Tools for Computational Dynamics in Industry and Academia**
Charlotte, North Carolina :: August 21 –24, 2016

Motivation



The Lagrangian-Lagrangian framework

- Based on the work behind Chrono::FSI 
- Fluid
 - Smoothed Particle Hydrodynamics (SPH)
- Solid
 - 3D rigid body dynamics (CM position, rigid rotation)
 - Absolute Nodal Coordinate Formulation (ANCF) for flexible bodies (nodes location and slope)
- Lagrangian-Lagrangian approach attractive since:
 - Consistent with Lagrangian tracking of discrete solid components
 - Straightforward simulation of free surface flows prevalent in target applications
 - Maps well to parallel computing architectures (GPU, many-core, distributed memory)
- *A Lagrangian-Lagrangian Framework for the Simulation of Fluid-Solid Interaction Problems with Rigid and Flexible Components, University of Wisconsin-Madison, 2014*



Smoothed Particle Hydrodynamics (SPH) method

- “Smoothed” refers to

$$\begin{aligned}
 f(\mathbf{x}) &= \int_S f(\mathbf{x}') \delta(\mathbf{x} - \mathbf{x}') d\mathbb{V} \\
 &= \int_S f(\mathbf{x}') W(\mathbf{x} - \mathbf{x}', h) d\mathbb{V} + O(h^2) \\
 &= \langle f(\mathbf{x}) \rangle + O(h^2)
 \end{aligned}$$


- “Particle” refers to

$$\begin{aligned}
 f(\mathbf{x}) &= \int_S \frac{f(\mathbf{x}')}{\rho(\mathbf{x}')} W(\mathbf{x} - \mathbf{x}', h) \rho(\mathbf{x}') d\mathbb{V} \\
 &\simeq \sum_b \frac{m_b}{\rho_b} f(\mathbf{x}_b) W(\mathbf{x} - \mathbf{x}_b, h)
 \end{aligned}$$

- Cubic spline kernel (often used)

$$W(q, h) = \frac{1}{4\pi h^3} \begin{cases} (2 - q)^3 - 4(1 - q)^3, & 0 \leq q < 1 \\ (2 - q)^3, & 1 \leq q < 2 \\ 0, & \text{otherwise} \end{cases}$$

$$\text{where } q \triangleq \frac{\|\mathbf{r}\|}{h}$$



Kernel Properties

$$\lim_{h \rightarrow 0} W(\mathbf{r}, h) = \delta(\mathbf{r})$$

$$W(\mathbf{r}, h) = W(-\mathbf{r}, h)$$

$$\int_S W(\mathbf{r}, h) d\mathbb{V} = 1$$

$$\lim_{\mathbf{r} \rightarrow \infty} W(\mathbf{r}, h) = 0$$

SPH for fluid dynamics

- Continuity

$$\frac{d\rho}{dt} = -\rho \nabla \cdot \mathbf{v}$$

$$\rho \nabla \cdot \mathbf{v} = \frac{\nabla \cdot (\rho^{\sigma-1} \mathbf{v}) - \mathbf{v} \cdot \nabla \rho^{\sigma-1}}{\rho^{\sigma-2}}.$$

- Momentum

$$\frac{d\mathbf{v}}{dt} = -\frac{\nabla p}{\rho} + \frac{\mu}{\rho} \nabla^2 \mathbf{v} + \mathbf{f}$$

$$\frac{\nabla p}{\rho} = \frac{p}{\rho^\sigma} \nabla \left(\frac{1}{\rho^{1-\sigma}} \right) + \rho^{\sigma-2} \nabla \left(\frac{p}{\rho^{\sigma-1}} \right)$$

- In the context of fluid dynamics, each particle carries fluid properties like pressure, density, etc.

$$\frac{d\rho_a}{dt} = \sum_b m_b \left(\frac{\mathbf{v}_a - \mathbf{v}_b}{\rho_a^{\sigma-2} \rho_b^{2-\sigma}} \right) \cdot \nabla_a W_{ab}$$

$$\frac{d\mathbf{v}}{dt} = -\sum_b m_b \left(\frac{p_a}{\rho_a^\sigma \rho_b^{2-\sigma}} + \frac{p_b}{\rho_a^{2-\sigma} \rho_b^\sigma} \right) \cdot \nabla_a W_{ab} + \sum_b m_b \frac{(\mu_a + \mu_b) \mathbf{x}_{ab} \cdot \nabla_a W_{ab}}{\bar{\rho}_{ab}^2 (x_{ab}^2 + \varepsilon \bar{h}_{ab}^2)} \mathbf{v}_{ab} + \mathbf{f}$$

$$\begin{aligned} \mathbf{x}_{ab} &= \mathbf{x}_a - \mathbf{x}_b \\ W_{ab} &= W(\mathbf{x}_{ab}, h) \\ \nabla_a &= \partial / \partial \mathbf{x}_a \end{aligned}$$


- Note: The above sums are done for millions of particles.

Fluid-Solid Interaction (ongoing work)


Boundary Condition Enforcing (BCE) markers for no-slip condition

- Rigidly attached to the solid body (hence their velocities are those of the corresponding material points on the solid)
- Hydrodynamic properties from the fluid



Rigid bodies/walls



Flexible Bodies



Example Representation





Current SPH Model

- Runge-Kutta 2nd order
 - Requires force calculation to happen twice per step
- Wall Boundary
 - Density changes for boundary particles as you would for the fluid particles.

$$\frac{d\rho_a}{dt} = \rho_a \sum_b \frac{m_b}{\rho_b} (\mathbf{v}_a - \mathbf{v}_b) \cdot \nabla_a W_{ab} \longleftrightarrow \rho_a = \sum_b m_b W_{ab}$$

- Periodic Boundary Condition
 - Markers who exit the periodic boundary, enter from the other side



Challenges for Scalable Distributed Memory Codes

- SPH is a computationally expensive method, hence, high performance computing (HPC) is necessary.
- High Performance Computing is hard.
 - MPI codes are able to achieve good strong and weak scaling, but... the developer is in charge of making this happen.
- Distributed memory challenges:
 - Communication bottlenecks > Computation bottlenecks
 - Load imbalance
 - Heterogeneity: processor types, process variation, memory hierarchies, etc.
 - Power/Temperature (becoming an important)
 - Fault tolerance
- To deal with these, we would like to seek
 - Not full automation
 - Not full burden on app-developers
 - But: a good division of labor between the system and app developers

Solution: Charm++

- Charm++ is a generalized approach to writing parallel programs
 - An alternative to the likes of MPI, UPC, GA etc.
 - But not to sequential languages such as C, C++, and Fortran
- Represents:
 - The style of writing parallel programs
 - The runtime system
 - And the entire ecosystem that surrounds it
- Three design principles:
 - **Overdecomposition, Migratability, Asynchrony**

Charm++ Design Principles

Overdecomposition

- Decompose work and data units into many more pieces than processing elements (cores, nodes, ...).
- Not so hard: problem decomposition needs to be done anyway.

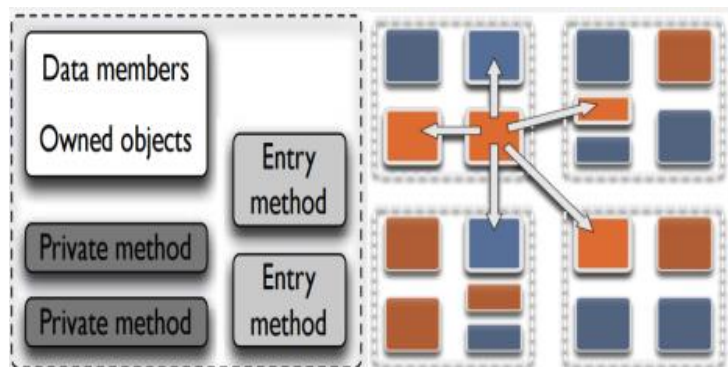



Figure 1: Single Charm Object (left). Overdecomposition; multiple chares in each execution unit exchanging data (right).

Migratability

- Allow data/work units to be migratable (by runtime and programmer).
- Communication is addressed to logical units (C++ objects) as opposed to physical units.
- Runtime System must keep track of these units



(b) Programmer's view: Collection of interacting chares

Asynchrony

- Message-driven execution
 - Let the work unit that happens to have data ("message") available execute next.
 - Runtime selects which work unit executes next (user can influence) → Scheduling






Figure 3: Compute idle time in MPI (top). Reduced idle times due to overdecomposition (bottom).

Realization of the design principle in Charm++

- Overdecomposed entities: **chares**
 - Chares are C++ objects
 - With methods designated as “entry” methods
 - Which can be invoked asynchronously by remote chares
 - Chares are organized into indexed collections
 - Each collection may have its own indexing scheme
 - 1D, ..7D
 - Sparse
 - Bitvector or string as an index
 - Chares communicate via asynchronous method invocations: **entry methods**
 - `A[i].foo(...)`; A is the name of a collection, i is the index of the particular chare.
- It is a kind of task-based parallelism
 - Pool of tasks + pool of workers
 - Runtime system selects what executes next.

Charm-based Parallel Model for SPH

- Hybrid decomposition (domain + force)
 - Inspired by NaMD (molecular dynamics application)
 - Domain Decomposition: 3D Cell Chare Array.
 - Each cell contains fluid/boundary/solid particles.
 - Data Units
 - Indexed: (x, y, z)
 - Force decomposition: 6D Compute Chare Array
 - Each compute chare is associated to a pair of cells.
 - Work units.
 - Indexed $(x_1, y_1, z_1, x_2, y_2, z_2)$
- No need to sort particles to find neighbor particles (overdecomposition implicitly takes care of it).
- Similar decomposition to LeanMD.
 - Charm++ Molecular Dynamics mini-app.
 - Kale, et al. "Charm++ for productivity and performance". PPL Technical Report, 2011.



Algorithm (Charm-based SPH)

1. Init each Cell Chare (very small subdomains)
2. For each subdomain create the number of Compute Chares

The following instructions happen in parallel for each Cell/Compute Chare.

Cell Array Loop (For each time step)	Compute Array Loop (For each time step)
3. SendPositions to each associate compute chare	4. When calcForces → SelfInteract OR Interact
6. Reduce forces from each compute chare	5. Send resulting forces
7. When reduce forces update properties at halfStep	
Repeat 3-7, but calc forces with marker properties at half step.	
8. Migrate Particles to Neighbors	
9. Load Balance every n steps	

Charm-based Parallel Model for FSI (ongoing work)

- Particles representing the solid will be contained with the fluid and boundary particles.
- Solid Chare Array (1D Array)
 - Particles keep track of the index of the solid they are associated with.
 - Once computes are done they send a message (invoke an entry method) to each solid they have particles of.
 - Do a force reduction and calculate the dynamics of the solid.

Charm++ In Practice

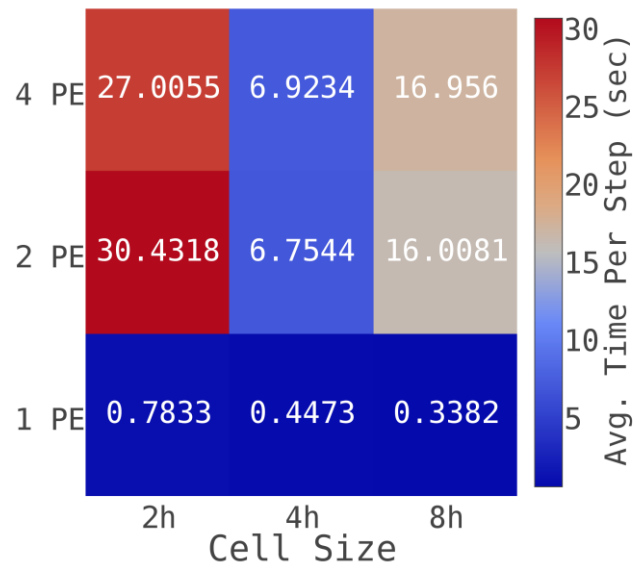
- Achieving optimal decomposition granularity
 - Average number of markers allowed per subdomain = Amount of work per chare.
 - Make sure there is enough work to hide communications.
 - Way too many chare objects is not optimal → Memory + Scheduling overheads
- Hyper Parameter Search
 - Vary Cell Size → Changes total number of cells and computes.
 - Vary Charm++ nodes per physical node → Feed comm network at max rate.
 - Varies number of communication and scheduling threads per node.
 - System specific. Small clusters might only need a single Charm++ node (1 communication thread), but larger clusters with different configurations might need more)

Charm++ Nodes\CellSize	2 * h	4 * h	8 * h
aprun -n 8 -N 1 -d 32 ./charmsph +ppn 31 +commap 0 +pemap 1-31	Average times per time step		
aprun -n 16 -N 2 -d 16 ./charmsph +ppn 15 +commap 0,16 +pemap 1-15:17-31			
aprun -n 32 -N 4 -d 8 ./charmsph +ppn 7 +commap 0,8,16,24 +pemap 1-7:9-15:17-23:25-31			

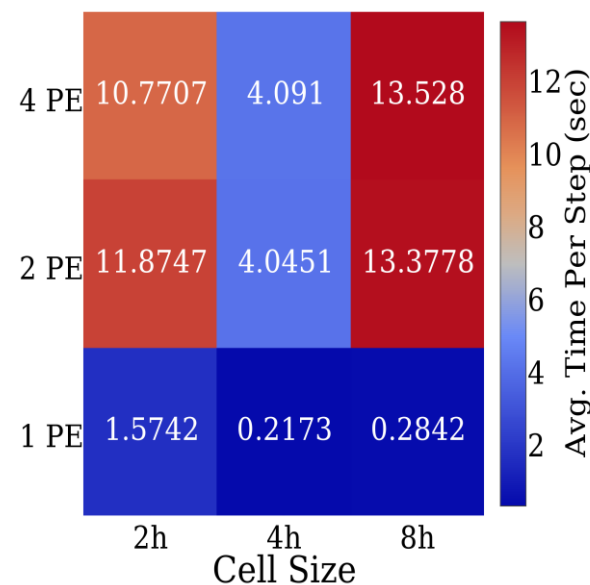
Results: Hyper parameter Search

- Hyper parameter search for optimal cell size and Charm++ nodes per physical node. Nodes denotes physical nodes (64 processors per node), and h denotes the particle interaction radius.
- H = Interaction radius of SPH particles.
- PE = Charm++ node (equivalent to MPI rank).

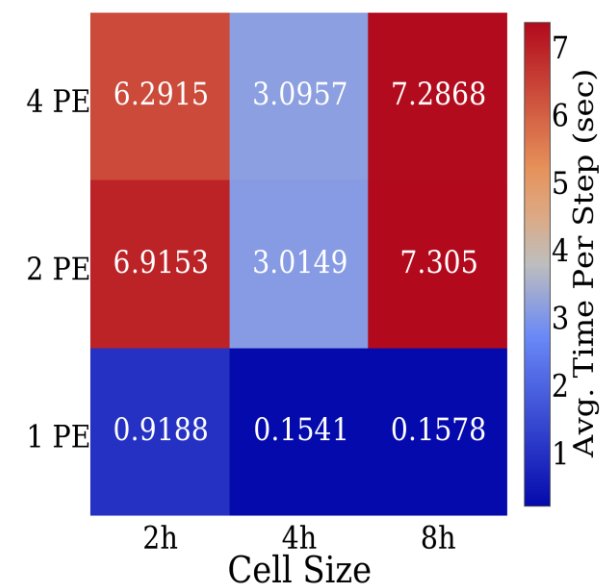
Nodes: 2, h: 0.025



Nodes: 4, h: 0.025



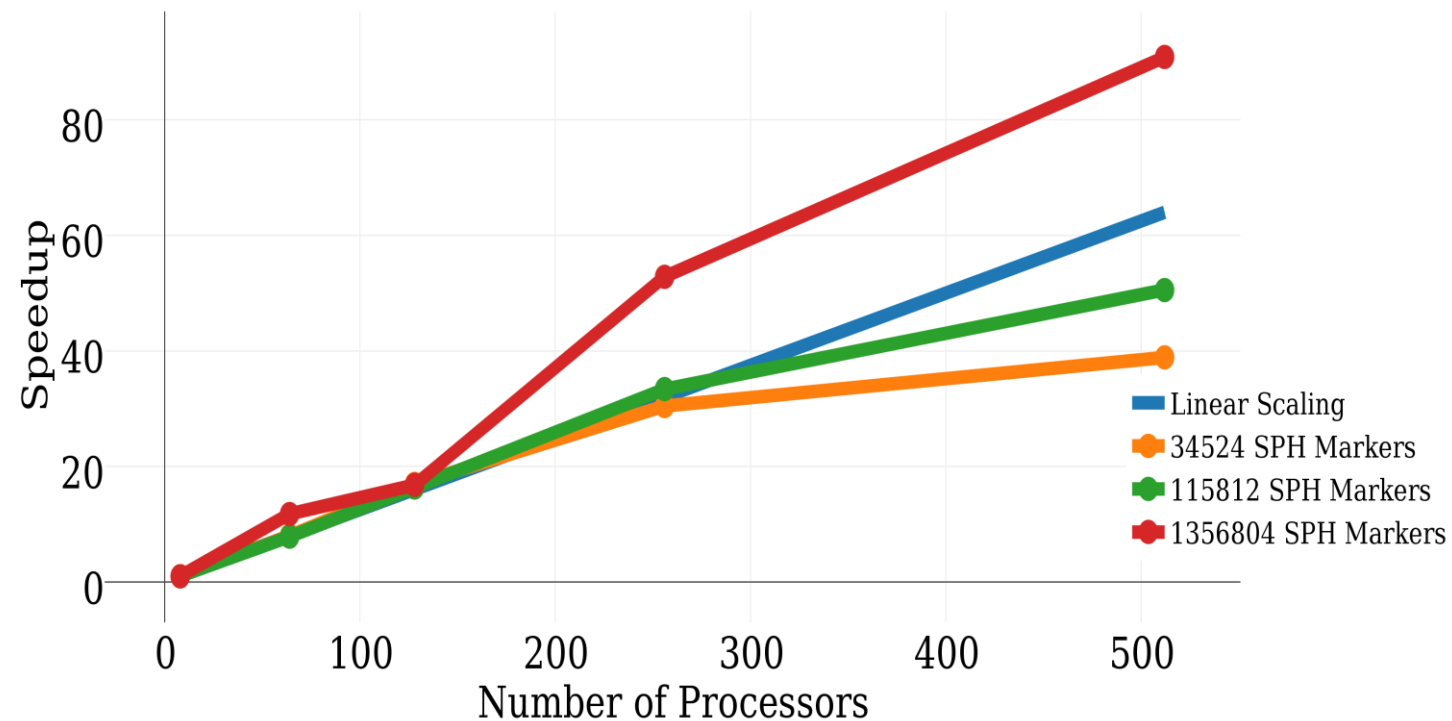
Nodes: 8, h: 0.025



Results: Strong Scaling

- Speedups calculated with respect to an 8 core run (8-504 cores).

Scalability with Optimal Parameters



Results: Dam break Simulation

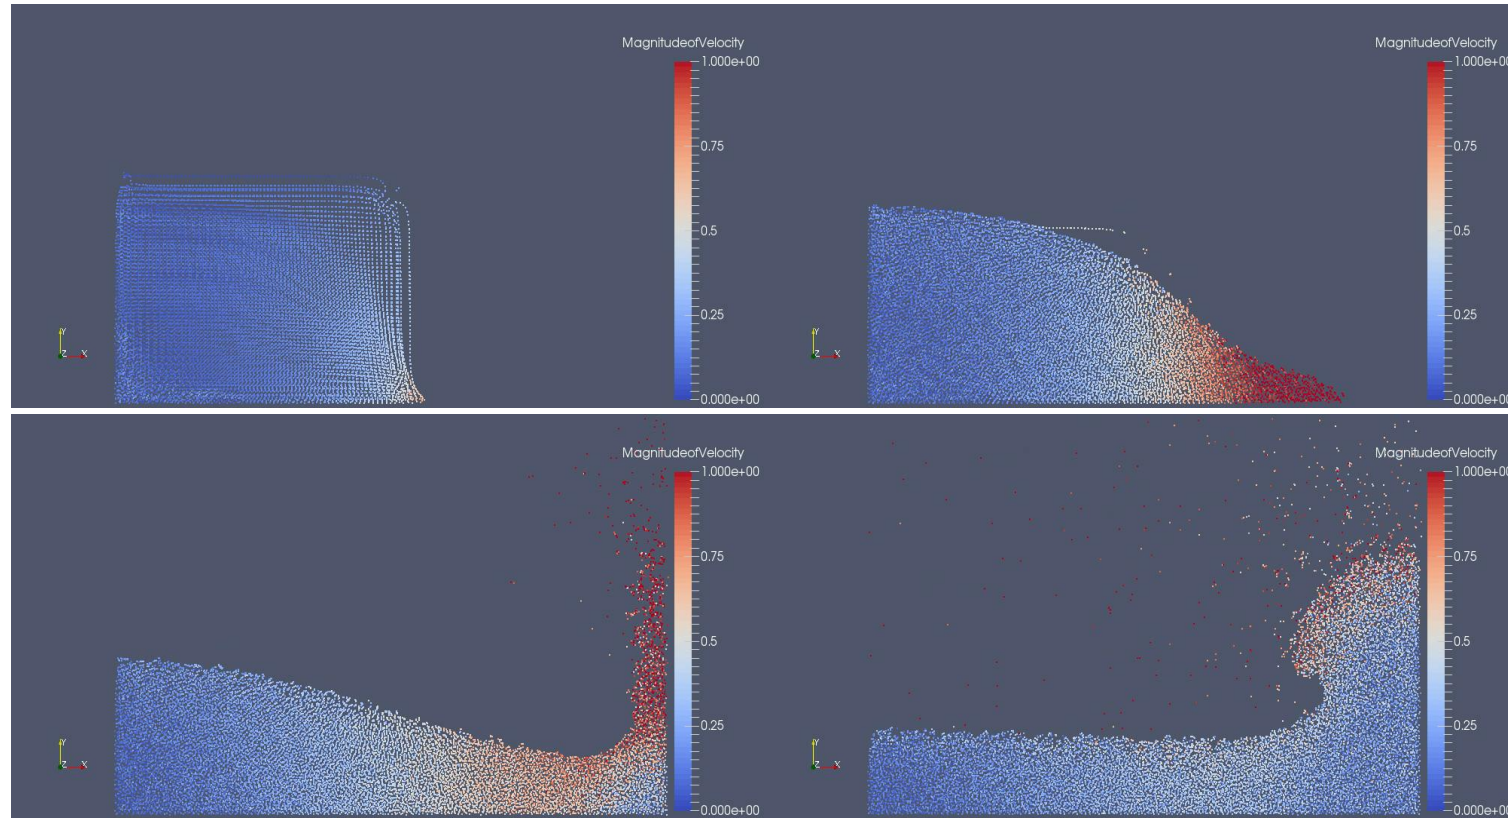


Figure 3: Dam break simulation (139,332 SPH Markers).

Note: Plain SPH requires hand tuning for stability.

Future Work (a lot to do)

- Improve the current SPH model following the same communication patterns for kernel calculations
 - Density Re-initialization.
 - Generalized Wall Boundary Condition
 - Adami, S., X. Y. Hu, and N. A. Adams. "A generalized wall boundary condition for smoothed particle hydrodynamics." *Journal of Computational Physics* 231.21 (2012): 7057-7075.
 - Pazouki, A., B. Song, and D. Negrut. "Technical Report TR-2015-09." (2015).
- Validation
- Hyper parameter search and scaling results on larger clusters.
 - Some bugs in HPC codes only appear after 1,000+ or 10,000+ cores.
- Performance+scaling comparison against other distributed memory SPH codes.
- Fluid-Solid Interaction
 - A. Pazouki, R. Serban, and D. Negrut, A Lagrangian-Lagrangian framework for the simulation of rigid and deformable bodies in fluid, *Multibody Dynamics: Computational Methods and Applications*, ISBN: 9783319072593, Springer, 2014.

Thank you!

Questions?

Code available at: <https://github.com/uwsbel/CharmSPH>