

RapidPT: A MATLAB toolbox for fast and efficient Permutation Tests

Felipe Gutierrez Barragan

<http://felipegb94.github.io/RapidPT/>

The Project - Overview


- **RapidPT** is a MATLAB toolbox for fast, scalable, and reliable permutation testing.
 - **Fast/Scalable:** Speedups of 20x – 80x when compared to a simple permutation testing implementation. 2x – 6x when compared to the state of the art, Statistical NonParametric Mapping (SnPM).
 - **Reliable:** Hundreds of validations runs against SnPM on neuroimaging datasets ranging from 20-400 subjects.
 - **Easy to use:** No setup, unlike hardware-based procedures. Will run fast on any modern machine.

Hypothesis Testing

- **Idea:** Calculate the probability that your claim/hypothesis is true.
- **Question:** Does the data display any interesting activity?
 - In neuroimaging – Voxel-wise differences between rest vs. activation? Healthy vs. Diseased?
- **Two Possibilities:**
 - Null Hypothesis vs. Alternate Hypothesis
 - H_0 vs. H_a

Hypothesis Testing - Procedure

1. Choose appropriate test statistic (t-test, mean difference, etc).
2. State H_0 and H_a
3. Construct the null distribution for the test statistic. This is the distribution of the statistic given that H_0 is true. This can be done analytically in some cases.
4. Compute the t-statistic with the given data.
5. Calculate the probability of observing such a statistic given that H_0 is true (p-value).
6. Accept or reject H_0 .



Hypothesis Testing - Types of Error

		Actual Situation “Truth”	
Decision		H_0 True	H_0 False
Do Not Reject H_0	Do Not Reject H_0	Correct Decision $1 - \alpha$	Incorrect Decision Type II Error β
	Reject H_0	Incorrect Decision Type I Error α	Correct Decision $1 - \beta$

$$\alpha = P(\text{Type I Error}) \quad \beta = P(\text{Type II Error})$$

Multiple Testing Problem

- **FWER:** Probability of making at least Type I Error (False Positive).

$$P(\text{Making at least 1 error in } m \text{ tests}) = 1 - (1 - \alpha)$$

$$P(\text{Making at least 1 error in } m \text{ tests}) = 1 - (1 - \alpha)^m$$

- For $m = 100 \rightarrow \text{FWER} = 0.99$.
- For a functional neuroimaging dataset we would do $>100,000$ tests. One for every voxel.
- Hence, we want to somehow control for the FWER and keep it under a certain probability α_0 .

Controlling FWER

- **Parametric Methods**

- Assumptions about the data and its distribution.
- **Bonferroni** – Conservative. Simply set $\alpha_0 = \alpha/v$ ($v = \text{number of tests}$).
- **Random Fields Theory** – Estimate number of activation areas, effectively reducing the number of tests.

- **Nonparametric/Resampling Methods**

- **Permutation Testing** – Empirically estimate the null distribution of the test statistics.
 - Exact control of the FWER.

Permutation Testing

- **Idea:** If the two groups do not differ, then I can permute the group/class labels and end up with approximately same set of t statistics.
- **Procedure:**
 1. Re-label images (permute the labels of the images).
 2. Compute test statistic for each voxel
 3. Repeat N times.
- After the procedure is done we will have the exact null distribution for each voxel, and we can proceed from step 4 of hypothesis testing .

Permutation Testing – Example, Single Test


	<u>Control</u>				<u>Drug</u>			
Expression	9	12	14	17	18	21	23	26
Average	13				22			

8 choose 4 = 70 possible permutations

Rearrangement of data									
Random Assignment	Control				Drug				Difference in Averages
1	9	12	14	17	18	21	23	26	9.0
2	9	12	14	18	17	21	23	26	8.5
3	9	12	14	21	17	18	23	26	7.0
4	9	12	14	23	17	18	21	26	6.0
5	9	12	14	26	17	18	21	23	4.5
6	9	12	17	18	14	21	23	26	7.0
7	9	12	17	21	14	18	23	26	5.5
8	9	12	17	23	14	18	21	26	4.5
9	9	12	17	26	14	18	21	23	3.0
10	9	12	18	21	14	17	23	26	5.0
11	9	12	18	23	14	17	21	26	4.0
12	9	12	18	26	14	17	21	23	2.5
13	9	12	21	23	14	17	18	26	2.5
14	9	12	21	26	14	17	18	23	1.0
15	9	12	23	26	14	17	18	21	0.0
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
69	18	21	23	26	9	12	14	17	-8.5
70	18	21	23	26	9	12	14	17	-9.0


14

Distribution of Difference between Treatment Means
Assuming No Treatment Effect



Computational Issues with Permutation Tests

Memory Issues



- Permutation Testing Matrix is very large
 - Let v be the number of voxels (around $5e5$)
 - Let T be the number of permutations (10,000)
 - Let P be the permutation testing matrix ($v \times T$)
 - Memory for that matrix is around **40 gigabytes**
 - *Note* – MATLAB crashes at some point when dealing with matrices that big.


Motivation: We need a lot of permutations...

- 1 test
 - 10,000 permutations → The smallest possible p-value is $\frac{1}{10,000}$.
 - Therefore the uncertainty near alpha = 0.01 is ± 0.0001 OR $\pm 1\%$
- 2 tests
 - 10,000 permutations → The smallest possible p-value is $\frac{1}{10,000}$ in each test.
 - Now we have 2 opportunities for a type I error
 - Therefore the uncertainty near alpha 0.01 in the worst case is ± 0.0002 or 2%
- 500,000 tests
 - 10,000 permutations → The smallest possible p-value is $\frac{1}{10,000}$ in each test.
 - Therefore the uncertainty in the worst case scenario is 100%


Idea: Look at the structure of P

- P is “highly structured” – A combination of low-rank signal and high-rank residual
- Example: MRI data 100 healthy vs. non-healthy. $v = 1,000$, $T = 2,000$

Permutation Testing Matrix P



Singular values of P



Core of RapidPT

- Many columns in P look similar to other columns as well as many rows look similar to other rows.
- If we compute a small number of entries of P we should be able to estimate the rest of it.
- Mathematically,

$$P = UW + S$$

- U is the low-rank basis of P .
 - W is the coefficient matrix
 - S is a high-rank random residual (noise).
- How many entries? In our experiments, subsampling $<1\%$ was enough

RapidPT Algorithm

Algorithm 3 Efficient Permutation Testing Algorithm

```
procedure RAPIDPT(data, nGroup1, numPermutations, numTrainingSamples, samplingRate)
   $\triangleright$  Training
2:    $P\_exact \leftarrow \text{PermTest}(\text{data}, nGroup1, numTrainingSamples)$ 
      $U \leftarrow \text{EstimateBasis}(P\_exact)$ 
4:    $W \leftarrow \text{EstimateCoeffMat}(U, P\_exact)$ 
      $S \leftarrow \text{ApproximateNoise}(U, P\_exact)$ 
   $\triangleright$  Recovery
6:    $maxnull \leftarrow \text{zeros}(1, numPermutations)$ 
      $subset \leftarrow \text{zeros}(1, samplingRate * numVoxels)$ 
8:   for  $i \leftarrow 1, numPermutations$  do
      $subset \leftarrow \text{randsample}(1 : numVoxels, samplingRate * numVoxels)$ 
10:     $p \leftarrow \text{PermTest}(\text{Data}(:, subset), nGroup1, 1)$ 
      $w \leftarrow \text{EstimateCoeffMat}(U(:, subU), p)$ 
12:     $s \leftarrow \text{EstimateNoise}(U(:, subU), p)$ 
      $p \leftarrow U * w + s$ 
14:     $maxnull(1, i) \leftarrow \max(p)$ 
     end for
16:   return  $maxnull$ 
end procedure
```

\triangleright Maximum Null Distribution

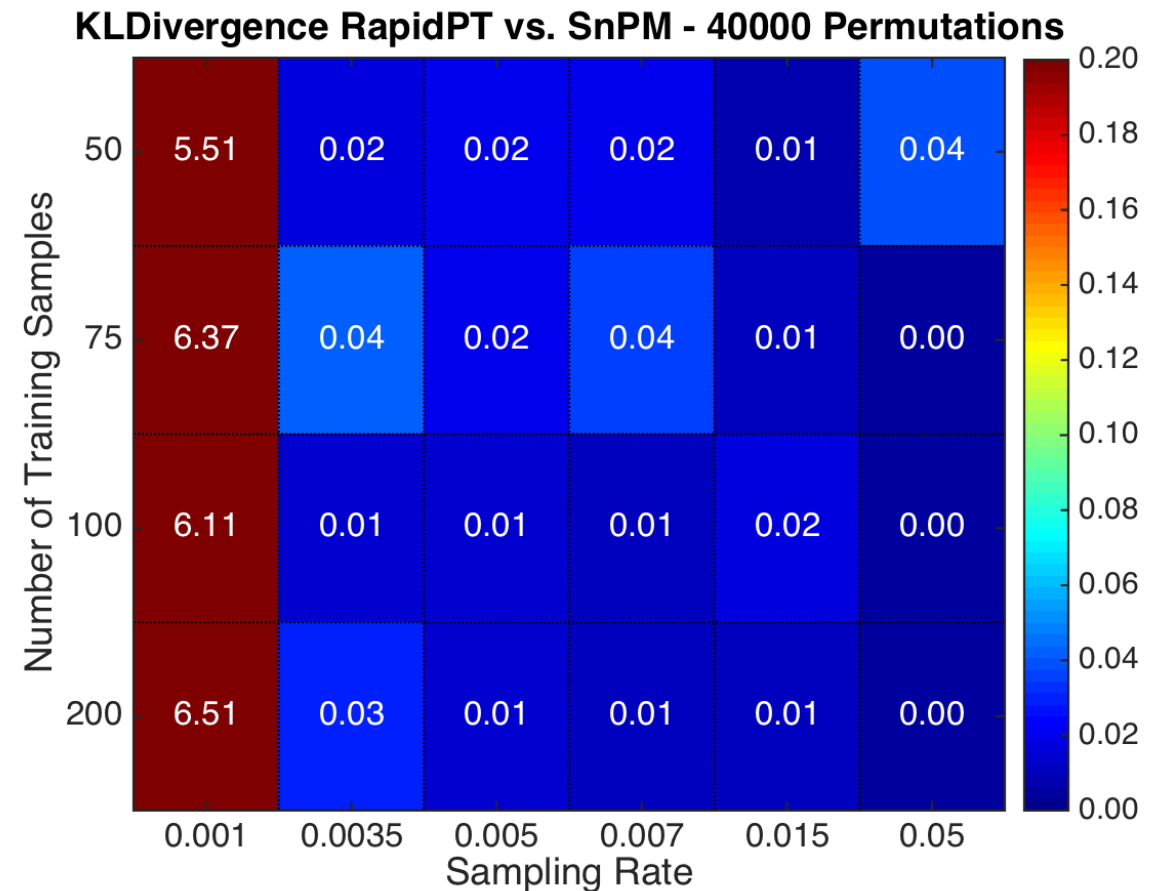
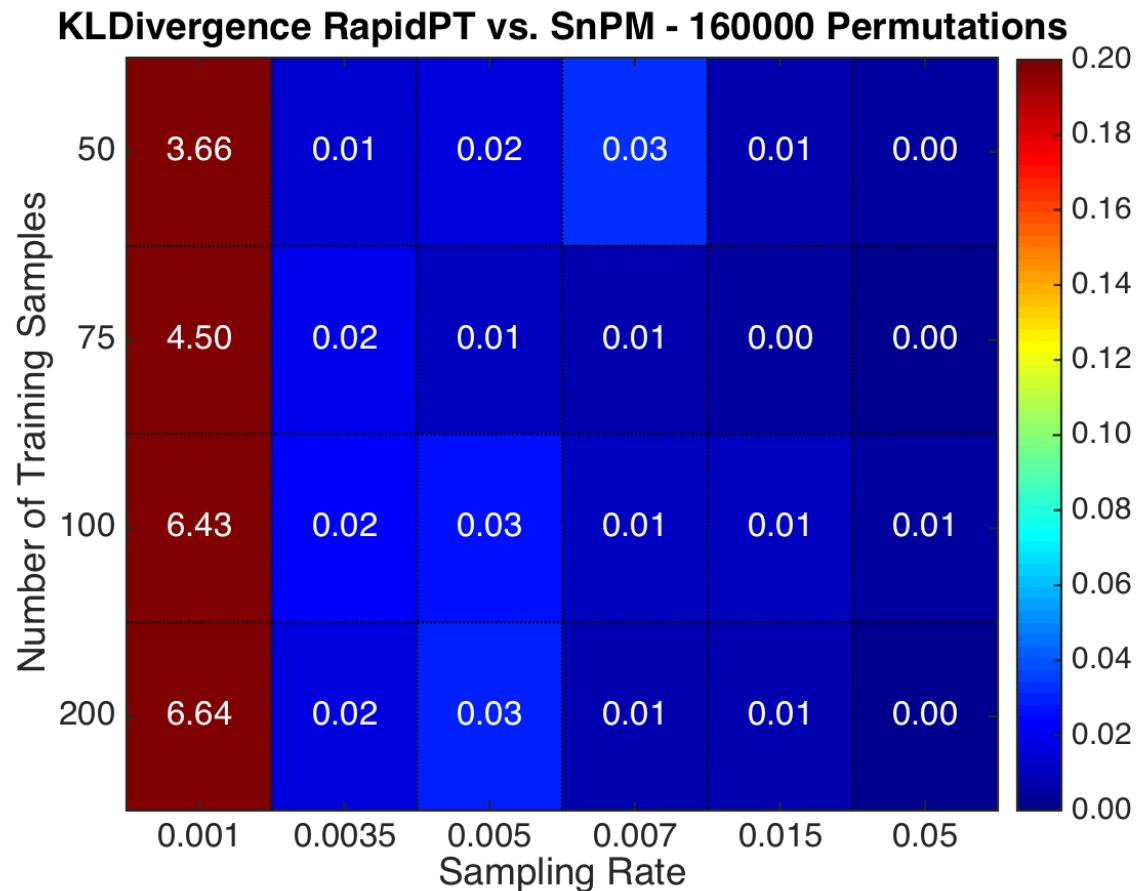
Evaluations Setup

- **Data:** T1 MRIs from ADNI2 are used.
 - 601 subjects (259 AD and 342 CN)
 - SPM preprocessing is applied.
 - GM images with approx. 500,000+ voxels are extracted.
 - Multiple combinations of dataset sizes
- **Experiments:** Can we recover the Maxnull distribution?
 - Stability of hyperparameters – Sub-sampling rates, and training samples
 - Computational Speedups (RapidPT vs. SnPM, RapidPT vs. NaivePT)


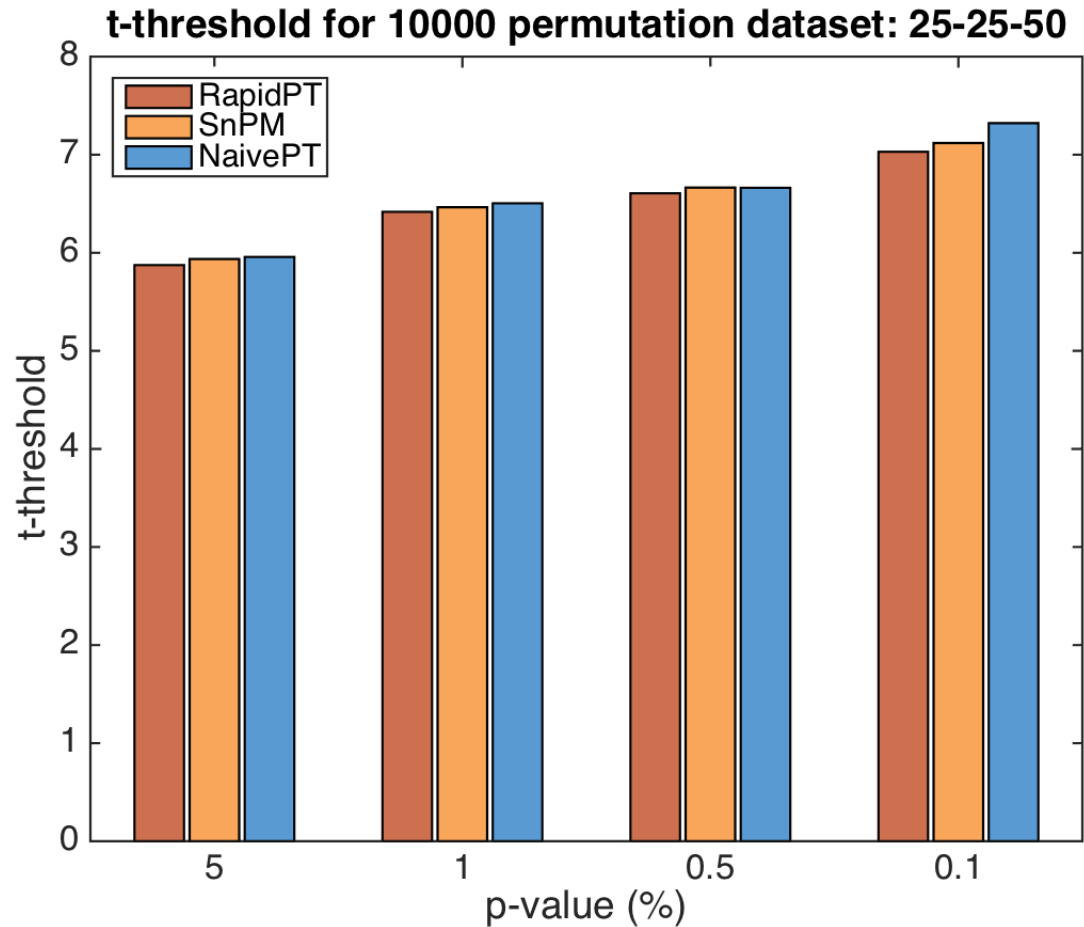
Recovered MaxNull Distribution Accuracy

KLDivergence: Measure of the difference between two distributions

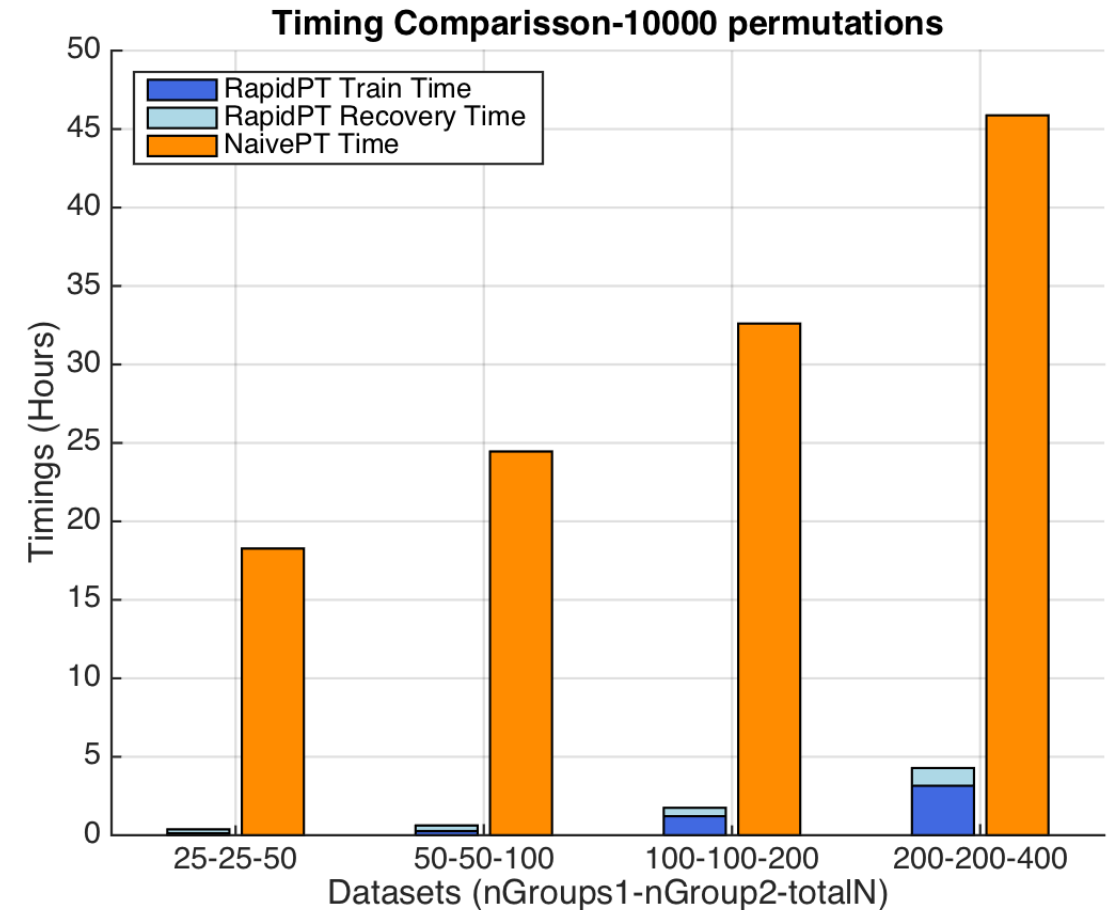

Dataset: 100 subject, 50 AD 50 CN



Calculated T-Threshold for a Given P-Value

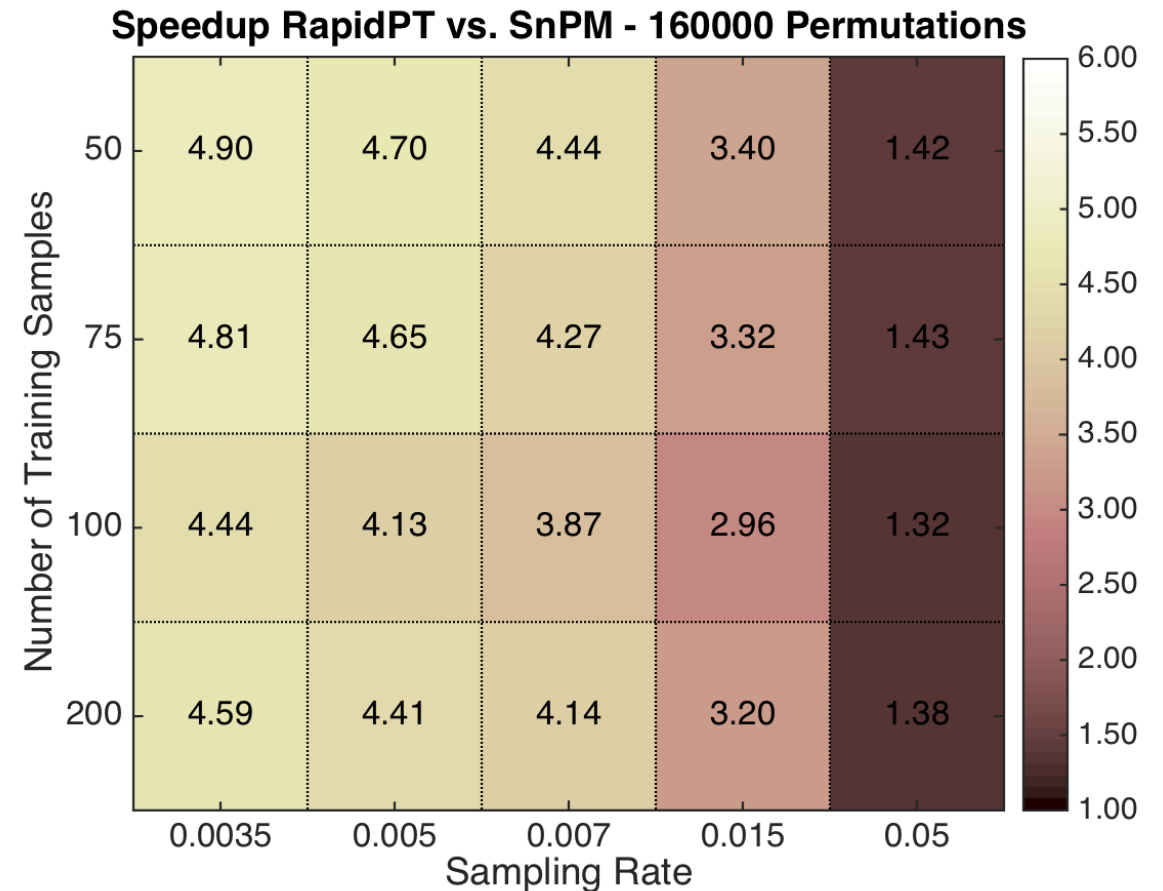



Speedup RapidPT vs. NaivePT



Speedup RapidPT vs. SnPM


Dataset used: 100 subjects, 50 AD 50 CN



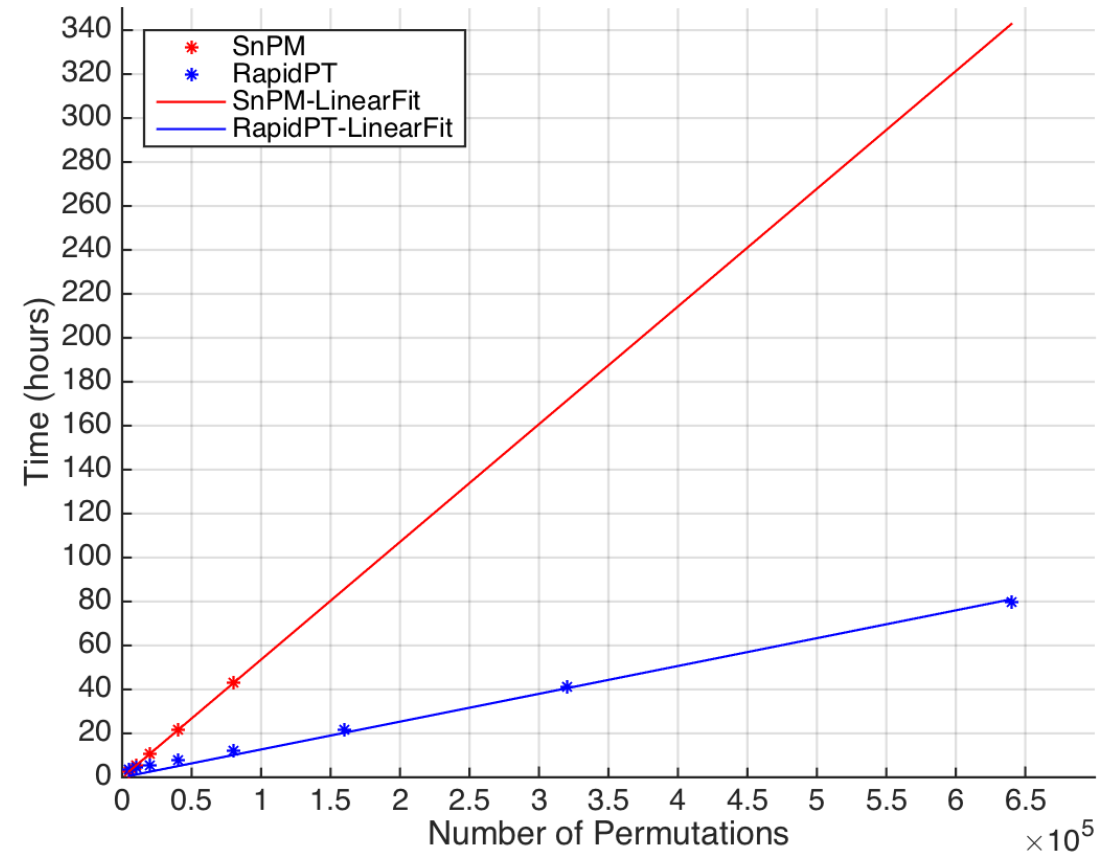
Scaling analysis RapidPT vs. SnPM

Putting the speedups into context...

SnPM vs RapidPT Performance-TwoSampleADRC-25-25-50

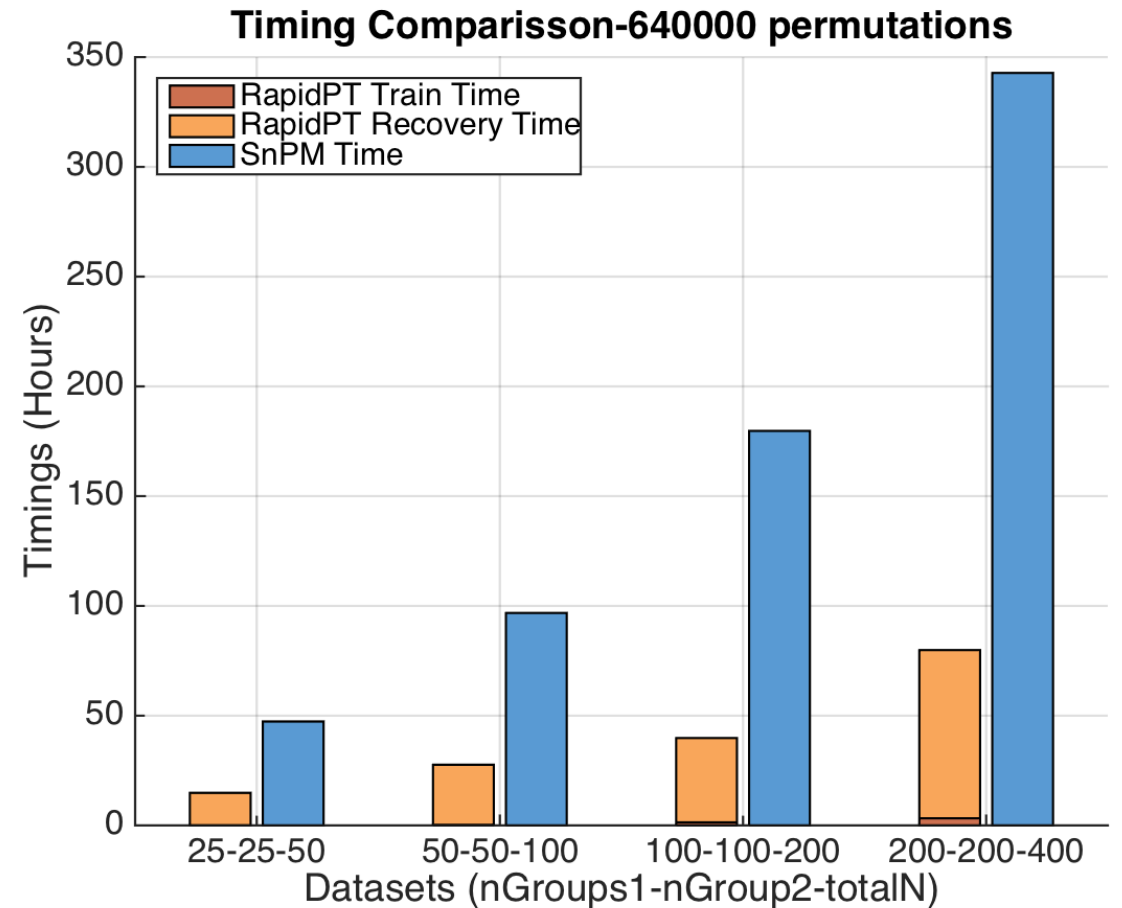



SnPM vs RapidPT Performance-TwoSampleADRC-200-200-400



Timings RapidPT vs. SnPM

Putting the speedups into context...



Discussion

- RapidPT achieves state of the art performance
 - It is able to recover the Maximum null distribution to a high degree of accuracy. Good control for FWER.
 - Fast – A few hours vs. days, and a day or two vs. a week or two..
 - Scalable – As the number of permutations increases and the dataset size increases the speedups against SnPM and NaivePT improve...
 - Easy to use????

RapidPT - Usage

Do pre-processing of .nii images for each subject and construct an NxV data matrix.


```
% Addpath RapidPT Repository Path (current working dir in this case)
RapidPTLibraryPath = '.';
addpath(RapidPTLibraryPath);

% Load input data and input labels
dataPath = '~/PermTest/data/ADRC/TwoSample/ADRC_50_25_25.mat';
% dataPath = '~/PermTest/data/face/Data_face.mat';
load(dataPath);
% N subjects, V voxels (or statistics)
[N,V] = size(Data);
numPermutations = 40000;
nGroup1 = 25; % You should what is the size of one of your groups prior.
% Set write to 1 if you want the matrices used to recover the permutation matrix.
% Setting this to 1 will make outputs a very large variable, but may be
% useful in certain cases.
write = 0;
|
[outputs, timings] = TwoSampleRapidPT(Data, numPermutations, nGroup1, write, RapidPTLibraryPath);

save(strcat('outputs/outputs_TwoSampleFace_',num2str(numPermutations),'.mat'),'outputs');
save(strcat('timings/timings_TwoSampleFace_',num2str(numPermutations),'.mat'),'timings');
```


RapidPT – Postprocess Example

```
% Get the outputs struct you obtained from RapidPT
load('~\PermTest/outputs/TwoSample_ADRC_200_200_400/rapidpt/outputs_80000_0
alpha = 0.01; % Significance level of 1 percent
tThresh_RapidPT = prctile(outputs.maxT, 100 - (100*alpha));
% Get the data
load('~\PermTest\data/ADRC/TwoSample/ADRC_400_200_200.mat');
[h,p,ci,stats]=ttest2(Data(1:200,:),Data(201:400,:),0.05,'both','unequal');
SampleMaxT = max(stats.tstat);
```



Warnings for RapidPT and Permutation Testing

- Small datasets and few permutations
 - A couple minutes vs one minute
 - Large uncertainty due to the smallest possible p-value being large.
 - 10 subjects \rightarrow 10 choose 5 = 252
 - Smallest possible p-val = $1/252$.



Acknowledgements and Website

- Vikas Singh
- Vamsi Ithapu
- ADNI and ADRC
- Repository and project website:
 - <https://github.com/felipegb94/RapidPT> (Repository)
 - <http://felipegb94.github.io/RapidPT/> (Website)

More Motivation: Implementation

- Implementation can easily be made inefficient
 - Inneficient data accesses
 - Not take advantage of MATLAB's optimized features

Algorithm 1 Traditional Two-Sample Permutation Testing. In this implementation `ttest2` refers to the two-sample t-test operation described in equation [1](#).

```
1: procedure TWOSAMPLENAIVEPT( $\mathbf{X}, r, N_1, N_2$ )
2:   for  $i \leftarrow 1, r$  do
3:     shuffleRows( $\mathbf{X}$ )
4:      $\mathbf{X}_1 \leftarrow \mathbf{X}(1 : N_1, :)$ 
5:      $\mathbf{X}_2 \leftarrow \mathbf{X}(N_1 + 1 : \text{end}, :)$ 
6:      $\mathbf{t} \leftarrow \text{ttest2}(\mathbf{X}_1, \mathbf{X}_2)$ 
7:      $\mathbf{T}(i, :) \leftarrow \mathbf{t}$ 
8:   end for
9: end procedure
```

▷ Mean and Variance Calculation

